

Reinforcement Learning

Terminologies

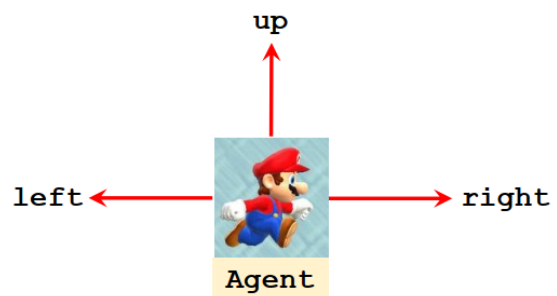
State and Action

Terminology: **state** and **action**

state s (this frame)



Action $a \in \{\text{left, right, up}\}$



state是场景

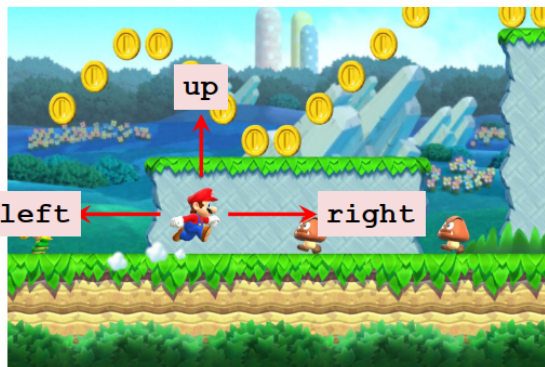
agent为智能体，也就是执行操作的对象，此处是马里奥，也可以是机器人，车等

action为agent做出的动作

Policy

Terminology: **policy**

policy π



- $\pi(a | s)$ is the probability of taking action $A = a$ given state s , e.g.,
 - $\pi(\text{left} | s) = 0.2$,
 - $\pi(\text{right} | s) = 0.1$,
 - $\pi(\text{up} | s) = 0.7$.
- Upon observing state $S = s$, the agent's action A can be random.

Policy函数，即策略函数，通常是一个概率分布函数，如该ppt中，在当前state下，马里奥选择往左走的策略的概率为0.2。该state下马里奥有三种策略，会在其中随机抽样，随机性使得policy更加灵活，难以被预测

Reward

Terminology: reward

reward R




- Collect a coin: $R = +1$
- Win the game: $R = +10000$
- Touch a Goomba: $R = -10000$ (game over).
- Nothing happens: $R = 0$

reward为奖励函数，如此处吃到金币+1分，通关+1w分，碰到敌人（game over）扣1w分

State Transition

Terminology: state transition

state transition



- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.
- $p(s'|s, a) = \mathbb{P}(S' = s' | S = s, A = a)$.

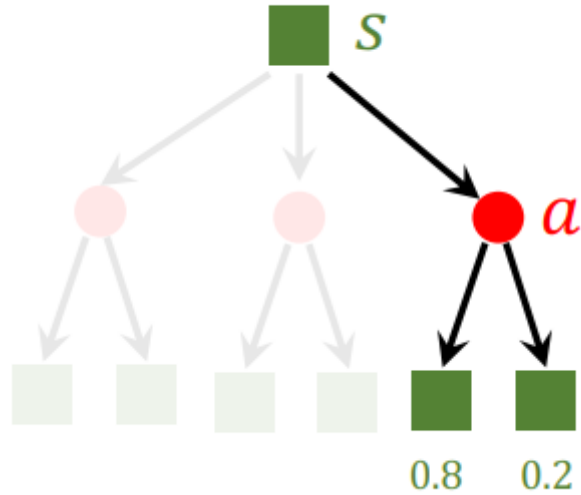
状态转移，顾名思义是old state变为new state，上图P为条件概率密度函数，意思是如果观测到当前的状态S以及动作A，下一个状态变成S一撇的概率

由于env以及action的随机性，所以状态转移具有随机性

Two Sources of Randomness

整个过程中的随机性主要来源于两点：

1. 是action的随机性，这个很好理解
2. 是state的随机性，在马里奥的例子中可以理解为敌人移动也是随机的，无法被agent知晓



- The randomness in **action** is from the policy function:

$$A \sim \pi(\cdot | s).$$

- The randomness in **state** is from the state-transition function:

$$S' \sim p(\cdot | s, a).$$

Trajectory

整个过程可以认为是


1. 观察state
2. 做出action
3. 观察到新的state并得到奖励（或惩罚）

Play game using AI

- (state, action, reward) trajectory:

$$\underline{s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_n, a_n, r_n.}$$

- One episode is from the beginning to the end (Mario wins or dies).



Rewards and Returns

returns 定义为未来所有 cumulative future reward 未来的累积奖励

$$U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$$

但其实之后的奖励对当前时刻不是同等重要的，如选择现在给你100元和一年后给你100元，大部分人会选择立刻得到一百

所以引入折扣回报 Discounted return

Definition: Discounted return (at time t).

$$U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^{n-t} R_n.$$

gamma 介于0到1，为超参数

R 与 S 和 A 有关

- Reward R_i depends on S_i and A_i .
- States can be random: $S_i \sim p(\cdot | s_{i-1}, a_{i-1})$.
- Actions can be random: $A_i \sim \pi(\cdot | s_i)$.
- If either S_i or A_i is random, then R_i is random.

所以U 与未来所有 S A 有关

At time t , the rewards, R_t, \dots, R_n , are random, so the return U_t is random.

- Reward R_i depends on S_i and A_i .
- U_t depends on R_t, R_{t+1}, \dots, R_n .
- $\rightarrow U_t$ depends on $S_t, A_t, S_{t+1}, A_{t+1}, \dots, S_n, A_n$.

Value Function

价值函数

Action-value function

U_t 其实是一个随机变量，他依赖于之后的所有动作和状态， t 时刻我们并不知道 U_t 是什么，所以我们可以对 U_t 求期望得到一个函数，即Action-value function，动作价值函数， Q_{π}

Definition: Action-value function for policy π .

$$Q_{\pi}(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t].$$

- Return U_t depends on actions $A_t, A_{t+1}, A_{t+2}, \dots$ and states $S_t, S_{t+1}, S_{t+2}, \dots$
- Actions are random: $\mathbb{P}[A = a | S = s] = \pi(a|s)$. (Policy function.)
- States are random: $\mathbb{P}[S' = s' | S = s, A = a] = p(s'|s, a)$. (State transition.)

此时我们 U_t 未知， S_t 和 A_t 是变量，且他们的概率分布已知（ S_t, A_t 分布），则可以用积分的方式把将来的SA对当前时刻 U_t 的影响通过积分求期望的方式获得

将 t 时刻之后的随机变量 A, S 都用积分积掉，之后得到的 Q_{π} 就只与当前时刻 t 的SA以及 π 有关

Action-value function的实际意义，一直policy函数 π 以及当前 t 时刻的 s ，则可以通过Action-value function Q_{π} 对每个action打分，看做出哪个action，最终 U_t 的期望最高

Optimal Action-value function

最优动作价值函数

之前说的Action-value function与 π, SA 有关，而我们有很多个policy函数 π ，我们要使得Action-value function最大，则可以做一个动态规划，取得最优的policy函数 π ，使得Action-value function最大，这样就可以消除 π 对Action-value function的影响（因为policy函数 π 已经确定了，不再是变量），得到一个最优的Action-value function，即Optimal Action-value function

Definition: Optimal action-value function.

$$Q^*(s_t, a_t) = \max_{\pi} Q_{\pi}(s_t, a_t).$$

Definition: State-value function.

$$\bullet V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)]$$

Definition: State-value function.

- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a)$. (Actions are discrete.)
- $V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \int \pi(a|s_t) \cdot Q_{\pi}(s_t, a) da$. (Actions are continuous.)

将 Q_{π} 对 A 求期望（将 A 当作随机变量），从而消掉 A

物理意义在于可以评估目前状态的胜算

Conclusion

Understanding the Value Functions

- **Action-value function:** $Q_{\pi}(s, a) = \mathbb{E} [U_t | S_t = s, A_t = a]$.
- Given policy π , $Q_{\pi}(s, a)$ evaluates how good it is for an agent to pick action a while being in state s .
- **State-value function:** $V_{\pi}(s) = \mathbb{E}_A [Q_{\pi}(s, A)]$
- For fixed policy π , $V_{\pi}(s)$ evaluates how good the situation is in state s .
- $\mathbb{E}_S [V_{\pi}(S)]$ evaluates how good the policy π is.

How does AI control the agent

1. policy-based learning 策略学习：已知 π ， S ，可以求得每个 A 的概率，再随机抽样
2. value-based learning 价值学习：求Optimal Action-value function Q-star

How does AI control the agent?

Suppose we have a good policy $\pi(a|s)$.

- Upon observe the state s_t ,
- random sampling: $a_t \sim \pi(\cdot | s_t)$.

Suppose we know the optimal action-value function $Q^*(s, a)$.

- Upon observe the state s_t ,
- choose the **action** that maximizes the value: $a_t = \operatorname{argmax}_a Q^*(s_t, a)$.

Value-Based Reinforcement Learning

Deep Q-Network (DQN)

回顾上节课讲的**Optimal Action-value function** Q_{star} , Q_{star} 的作用是判断在当前state下, 哪个action带来的未来reward总和的期望越大。而 Q_{star} 往往是不能直接得到的, 价值学习的基本想法就是通过学习一个函数(神经网络)来近似 Q_{star}

Approximate the Q Function

Goal: Win the game (\approx maximize the total reward.)

Question: If we know $Q^*(s, a)$, what is the best **action**?

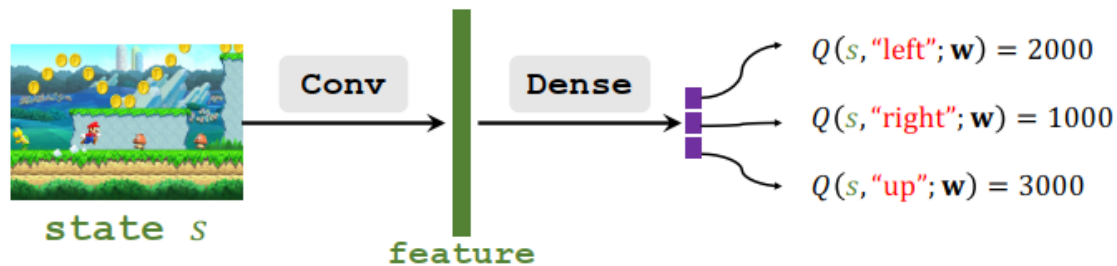
- Obviously, the best action is $a^* = \operatorname{argmax}_a Q^*(s, a)$.

Challenge: We do not know $Q^*(s, a)$.

- Solution: Deep Q Network (**DQN**)
- Use neural network $Q(s, a; \mathbf{w})$ to approximate $Q^*(s, a)$.

Deep Q Network (DQN)

- Input shape: size of the screenshot.
- Output shape: dimension of action space.



Question: Based on the predictions, what should be the **action**?

Temporal Difference (TD) Learning


王老师举了一个预测开车时间的例子

Example

- I want to drive from NYC to Atlanta.
- Model $Q(\mathbf{w})$ estimates the time cost, e.g., 1000 minutes.

Question: How do I update the model?


- Make a prediction: $q = Q(\mathbf{w})$, e.g., $q = 1000$.
- Finish the trip and get the target y , e.g., $y = 860$.
- Loss: $L = \frac{1}{2}(q - y)^2$.
- Gradient: $\frac{\partial L}{\partial \mathbf{w}} = \frac{\partial q}{\partial \mathbf{w}} \cdot \frac{\partial L}{\partial q} = (q - y) \cdot \frac{\partial Q(\mathbf{w})}{\partial \mathbf{w}}$.
- Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L}{\partial \mathbf{w}} \big|_{\mathbf{w}=\mathbf{w}_t}$.



如图用梯度下降法，比较naive，需要完成一次旅程才能update model

Temporal Difference (TD) Learning

- Model's estimate: $Q(\mathbf{w}) = 1000$ minutes.
- Updated estimate: $300 + 600 = 900$ minutes.
 ↘
 TD target.
- TD target $y = 900$ is a more reliable estimate than 1000 .
- Loss: $L = \frac{1}{2}(Q(\mathbf{w}) - y)^2$.
- Gradient: $\frac{\partial L}{\partial \mathbf{w}} = (1000 - 900) \cdot \frac{\partial Q(\mathbf{w})}{\partial \mathbf{w}}$.
- Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_t}$.



使用TD learning，走到路程中间（300min处），再使用模型预测一次，预测值为600min，容易想到离终点越接近，该估计会越准，所以可以认为300+600=900的估计比一开始的1000更准，1000与900的差称为TD error

How to apply TD learning to DQN?

- In the “driving time” example, we have the equation:

$$T_{\text{NYC} \rightarrow \text{ATL}} \approx T_{\text{NYC} \rightarrow \text{DC}} + T_{\text{DC} \rightarrow \text{ATL}}.$$

Model's estimate

Actual time

Model's estimate

- In deep reinforcement learning:

$$Q(s_t, a_t; \mathbf{w}) \approx r_t + \gamma \cdot Q(s_{t+1}, a_{t+1}; \mathbf{w}).$$

TD learning可以运用的场景，即可以写作estimate = estimate+actual的形式，其中的等于是我们最理想的情况，即TD error等于0

回顾discount return

Identity: $U_t = R_t + \gamma \cdot U_{t+1}.$

$$\begin{aligned}
 U_t &= R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \gamma^4 \cdot R_{t+4} + \dots \\
 &= R_t + \gamma \cdot (R_{t+1} + \gamma \cdot R_{t+2} + \gamma^2 \cdot R_{t+3} + \gamma^3 \cdot R_{t+4} + \dots) \\
 &\quad \underbrace{\hspace{10em}} \\
 &= U_{t+1}
 \end{aligned}$$

How to apply TD learning to DQN?

Identity: $U_t = R_t + \gamma \cdot U_{t+1}$.

TD learning for DQN:

- DQN's output, $Q(s_t, a_t; \mathbf{w})$, is an estimate of U_t .
- DQN's output, $Q(s_{t+1}, a_{t+1}; \mathbf{w})$, is an estimate of U_{t+1} .

• Thus,
$$\underbrace{Q(s_t, a_t; \mathbf{w})}_{\text{Prediction}} \approx \underbrace{r_t + \gamma \cdot Q(s_{t+1}, a_{t+1}; \mathbf{w})}_{\text{TD target}}.$$

Train DQN using TD learning

- Prediction: $Q(s_t, a_t; \mathbf{w}_t)$.
- TD target:

$$\begin{aligned} y_t &= r_t + \gamma \cdot Q(s_{t+1}, a_{t+1}; \mathbf{w}_t) \\ &= r_t + \gamma \cdot \max_a Q(s_{t+1}, a; \mathbf{w}_t). \end{aligned}$$

- Loss: $L_t = \frac{1}{2} [Q(s_t, a_t; \mathbf{w}) - y_t]^2$.
- Gradient descent: $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial L_t}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}_t}$.


Policy-Based Reinforcement Learning

同样，策略函数 π 也是难以直接获得的，所以需要通过神经网络来近似，此神经网络被称为policy network

- Use policy network $\pi(a|s; \boldsymbol{\theta})$ to approximate $\pi(a|s)$.

Policy Network $\pi(a|s; \theta)$

- $\sum_{a \in \mathcal{A}} \pi(a|s; \theta) = 1$.
- Here, $\mathcal{A} = \{\text{"left"}, \text{"right"}, \text{"up"}\}$ is the set all actions.
- That is why we use softmax activation.



回顾之前学的state-value fuction V_{π} , V_{π} 是 Q_{π} 对action求期望, 可以表示在当前state下的胜算

$$\bullet V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)]$$

V_{π} 可以写作下图形式

Definition: State-value function.

$$\bullet V_{\pi}(s_t) = \mathbb{E}_A [Q_{\pi}(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_{\pi}(s_t, a).$$

Approximate state-value function.

- Approximate policy function $\pi(a|s_t)$ by policy network $\pi(a|s_t; \theta)$.
- Approximate value function $V_{\pi}(s_t)$ by:

$$V(s_t; \theta) = \sum_a \pi(a|s_t; \theta) \cdot Q_{\pi}(s_t, a).$$

$V(s, \theta)$ 可以度量状态 s 和策略网络 θ 的好坏, 给定状态 s , 策略网络 θ 越好, 则 V 越大
所以我们可以把目标函数设置为

Policy-based learning: Learn θ that maximizes $J(\theta) = \mathbb{E}_s [V(s; \theta)]$.

策略网络 θ 越好, J_{θ} 越大

How to improve θ ? Policy gradient ascent!

- Observe state s .
- Update policy by: $\theta \leftarrow \theta + \beta \cdot \frac{\partial V(s; \theta)}{\partial \theta}$

Policy gradient

此处使用的是梯度上升算法, 因为我们是想要目标函数越大越好 (对比loss函数)

Policy gradient

此处Policy gradient的求导会用到高数和概率论，不太好记录，建议多看看ppt和视频

Policy Gradient

Definition: Approximate state-value function.

$$\bullet V(s; \theta) = \sum_a \pi(a|s; \theta) \cdot Q_\pi(s, a).$$

Policy gradient: Derivative of $V(s; \theta)$ w.r.t. θ .

$$\begin{aligned} \bullet \frac{\partial V(s; \theta)}{\partial \theta} &= \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \sum_a \pi(a|s; \theta) \cdot \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a) \\ &= \mathbb{E}_A \left[\frac{\partial \log \pi(A|s; \theta)}{\partial \theta} \cdot Q_\pi(s, A) \right]. \end{aligned}$$

The expectation is taken w.r.t. the random variable $A \sim \pi(\cdot | s; \theta)$.

Two forms of policy gradient:

- Form 1: $\frac{\partial V(s; \theta)}{\partial \theta} = \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a).$
- Form 2: $\frac{\partial V(s; \theta)}{\partial \theta} = \mathbb{E}_{A \sim \pi(\cdot | s; \theta)} \left[\frac{\partial \log \pi(A|s; \theta)}{\partial \theta} \cdot Q_\pi(s, A) \right].$

我们得到了以上两种方式来表示policy gradient

对于动作是离散形式，可以使用Form1枚举计算

Calculate Policy Gradient for Discrete Actions

If the actions are **discrete**, e.g., action space $\mathcal{A} = \{\text{"left"}, \text{"right"}, \text{"up"}\}, \dots$

Use **Form 1**: $\frac{\partial V(s; \theta)}{\partial \theta} = \sum_a \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a).$

1. Calculate $f(a, \theta) = \frac{\partial \pi(a|s; \theta)}{\partial \theta} \cdot Q_\pi(s, a)$, for every action $a \in \mathcal{A}$.
2. Policy gradient: $\frac{\partial V(s; \theta)}{\partial \theta} = f(\text{"left"}, \theta) + f(\text{"right"}, \theta) + f(\text{"up"}, \theta).$

This approach **does not work** for **continuous** actions.

对于action是连续形式，则需要积分，但Qpi是一个神经网络非常复杂，不能直接积分得到解析解，所以需要使用蒙特卡洛算法近似（此处需要补概率论...）

Calculate Policy Gradient

$$\text{Policy Gradient: } \frac{\partial V(s; \theta)}{\partial \theta} = \mathbb{E}_{A \sim \pi(\cdot | s; \theta)} \left[\frac{\partial \log \pi(A | s, \theta)}{\partial \theta} \cdot Q_{\pi}(s, A) \right].$$

1. Randomly sample an action \hat{a} according to $\pi(\cdot | s; \theta)$.
2. Calculate $\mathbf{g}(\hat{a}, \theta) = \frac{\partial \log \pi(\hat{a} | s; \theta)}{\partial \theta} \cdot Q_{\pi}(s, \hat{a})$.
3. Use $\mathbf{g}(\hat{a}, \theta)$ as an approximation to the policy gradient $\frac{\partial V(s; \theta)}{\partial \theta}$.

1. 随机抽样一个a_hat, 抽样是根据概率密度函数pi抽的
2. 计算g(a_hat, theta)

- By the definition of \mathbf{g} , $\mathbb{E}_A[\mathbf{g}(A, \theta)] = \frac{\partial V(s; \theta)}{\partial \theta}$.
- $\mathbf{g}(\hat{a}, \theta)$ is an unbiased estimate of $\frac{\partial V(s; \theta)}{\partial \theta}$.

可以知道g(A, theta) 对A求期望即为V的导数

且g(a_hat, theta) 是V求导的无偏估计 (?)

则可以用g(a_hat, theta) 来近似V求导

蒙特卡洛算法就是通过抽取一个或多个样本对期望进行近似

整个流程如下图所示

Algorithm

1. Observe the state s_t .
2. Randomly sample action a_t according to $\pi(\cdot | s_t; \theta_t)$.
3. Compute $q_t \approx Q_{\pi}(s_t, a_t)$ (some estimate).
4. Differentiate policy network: $\mathbf{d}_{\theta, t} = \frac{\partial \log \pi(a_t | s_t, \theta)}{\partial \theta} \big|_{\theta = \theta_t}$.
5. (Approximate) policy gradient: $\mathbf{g}(a_t, \theta_t) = q_t \cdot \mathbf{d}_{\theta, t}$.
6. Update policy network: $\theta_{t+1} = \theta_t + \beta \cdot \mathbf{g}(a_t, \theta_t)$.

但还有一个问题没有解决, 即由于Qpi无法得知, 所以qt不能直接得到, 有如下两种方法来近似

1. Reinforce算法

因为Qpi的Ut的期望, 所以可以用ut来近似Qpai, 即近似qt, 该方法需要完整玩完一局游戏才能对策略函数进行更新

2. Randomly sample action a_t according to $\pi(\cdot | s_t; \theta_t)$.

3. Compute $q_t \approx Q_\pi(s_t, a_t)$ (some estimate). **How?**

Option 1: REINFORCE.

- Play the game to the end and generate the trajectory:

$$s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T.$$

- Compute the discounted return $u_t = \sum_{k=t}^T \gamma^{k-t} r_k$, for all t .
- Since $Q_\pi(s_t, a_t) = \mathbb{E}[U_t]$, we can use u_t to approximate $Q_\pi(s_t, a_t)$.
- \Rightarrow Use $q_t = u_t$.

2. actor-critic method

用神经网络近似 q_t ，下节课具体讲解